

# Cyclistic bike-share analysis case study

Ramya Bandi

2024-11-30

```
library(tidyverse) #helps wrangle data
```

## Use the conflicted package to manage conflicts

```
library(conflicted)
```

## Set dplyr::filter and dplyr::lag as the default choices

```
conflict_prefer("filter", "dplyr")
```

```
conflict_prefer("lag", "dplyr")
```

```
#=====
```

## STEP 1: COLLECT DATA

```
#=====
```

## Upload Divvy datasets (csv files) here

```
q1_2019 <- read_csv("Divvy_Trips_2019_Q1.csv")
```

```
## Rows: 365069 Columns: 12
## -- Column specification -----
## Delimiter: ","
## chr (6): start_time, end_time, from_station_name, to_station_name, usertype,...
## dbl (5): trip_id, bikeid, from_station_id, to_station_id, birthyear
## num (1): tripduration
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
q1_2020 <- read_csv("Divvy_Trips_2020_Q1.csv")
```

```
## Rows: 426887 Columns: 13
## -- Column specification -----
## Delimiter: ","
## chr (7): ride_id, rideable_type, started_at, ended_at, start_station_name, e...
## dbl (6): start_station_id, end_station_id, start_lat, start_lng, end_lat, en...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

#####

## STEP 2: WRANGLE DATA AND COMBINE INTO A SINGLE FILE

#####

Compare column names each of the files

While the names don't have to be in the same order, they DO need to match perfectly before

we can use a command to join them into one file

```
colnames(q1_2019)
```

```
## [1] "trip_id"          "start_time"       "end_time"
## [4] "bikeid"           "tripduration"     "from_station_id"
## [7] "from_station_name" "to_station_id"    "to_station_name"
## [10] "usertype"         "gender"           "birthyear"
```

```
colnames(q1_2020)
```

```
## [1] "ride_id"          "rideable_type"    "started_at"
## [4] "ended_at"         "start_station_name" "start_station_id"
## [7] "end_station_name" "end_station_id"    "start_lat"
## [10] "start_lng"        "end_lat"          "end_lng"
## [13] "member_casual"
```

Rename columns to make them consistent with q1\_2020 (as this will be the supposed

going-forward table design for Divvy)

```
(q1_2019 <- rename(q1_2019
,ride_id = trip_id
,rideable_type = bikeid
,started_at = start_time
,ended_at = end_time
,start_station_name = from_station_name
,start_station_id = from_station_id
,end_station_name = to_station_name
,end_station_id = to_station_id
,member_casual = usertype
))
```

```
## # A tibble: 365,069 x 12
```

```
##   ride_id started_at ended_at rideable_type tripduration start_station_id
##   <dbl> <chr>      <chr>      <dbl>      <dbl>      <dbl>
## 1 21742443 01-01-2019 00:~ 01-01-2~      2167        390        199
## 2 21742444 01-01-2019 00:~ 01-01-2~      4386        441         44
## 3 21742445 01-01-2019 00:~ 01-01-2~      1524        829         15
## 4 21742446 01-01-2019 00:~ 01-01-2~       252       1783        123
```

```
## 5 21742447 01-01-2019 00:~ 01-01-2~          1170          364          173
## 6 21742448 01-01-2019 00:~ 01-01-2~          2437          216           98
## 7 21742449 01-01-2019 00:~ 01-01-2~          2708          177           98
## 8 21742450 01-01-2019 00:~ 01-01-2~          2796          100          211
## 9 21742451 01-01-2019 00:~ 01-01-2~          6205          1727          150
## 10 21742452 01-01-2019 00:~ 01-01-2~          3939          336          268
## # i 365,059 more rows
## # i 6 more variables: start_station_name <chr>, end_station_id <dbl>,
## #   end_station_name <chr>, member_casual <chr>, gender <chr>, birthyear <dbl>
```

## Inspect the dataframes and look for incongruencies

```
str(q1_2019)
```

```
## spc_tbl_ [365,069 x 12] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
##  $ ride_id          : num [1:365069] 21742443 21742444 21742445 21742446 21742447 ...
##  $ started_at       : chr [1:365069] "01-01-2019 00:04" "01-01-2019 00:08" "01-01-2019 00:13" "01-01-2019 00:17" ...
##  $ ended_at         : chr [1:365069] "01-01-2019 00:11" "01-01-2019 00:15" "01-01-2019 00:27" "01-01-2019 00:31" ...
##  $ rideable_type     : num [1:365069] 2167 4386 1524 252 1170 ...
##  $ tripduration      : num [1:365069] 390 441 829 1783 364 ...
##  $ start_station_id  : num [1:365069] 199 44 15 123 173 98 98 211 150 268 ...
##  $ start_station_name: chr [1:365069] "Wabash Ave & Grand Ave" "State St & Randolph St" "Racine Ave & Grand Ave" "Clark St & Leland Ave" ...
##  $ end_station_id    : num [1:365069] 84 624 644 176 35 49 49 142 148 141 ...
##  $ end_station_name  : chr [1:365069] "Milwaukee Ave & Grand Ave" "Dearborn St & Van Buren St (*)" "Clark St & Leland Ave" "Clark St & Leland Ave" ...
##  $ member_casual     : chr [1:365069] "Subscriber" "Subscriber" "Subscriber" "Subscriber" ...
##  $ gender            : chr [1:365069] "Male" "Female" "Female" "Male" ...
##  $ birthyear         : num [1:365069] 1989 1990 1994 1993 1994 ...
##  - attr(*, "spec")=
##    .. cols(
##      ..   trip_id = col_double(),
##      ..   start_time = col_character(),
##      ..   end_time = col_character(),
##      ..   bikeid = col_double(),
##      ..   tripduration = col_number(),
##      ..   from_station_id = col_double(),
##      ..   from_station_name = col_character(),
##      ..   to_station_id = col_double(),
##      ..   to_station_name = col_character(),
##      ..   usertype = col_character(),
##      ..   gender = col_character(),
##      ..   birthyear = col_double()
##    .. )
##  - attr(*, "problems")=<externalptr>
```

```
str(q1_2020)
```

```
## spc_tbl_ [426,887 x 13] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
##  $ ride_id          : chr [1:426887] "EACB19130B0CDA4A" "8FED874C809DC021" "789F3C21E472CA96" "C9A330B0CDA4A" ...
##  $ rideable_type     : chr [1:426887] "docked_bike" "docked_bike" "docked_bike" "docked_bike" ...
##  $ started_at       : chr [1:426887] "21-01-2020 20:06" "30-01-2020 14:22" "09-01-2020 19:29" "06-01-2020 19:29" ...
##  $ ended_at         : chr [1:426887] "21-01-2020 20:14" "30-01-2020 14:26" "09-01-2020 19:32" "06-01-2020 19:32" ...
##  $ start_station_name: chr [1:426887] "Western Ave & Leland Ave" "Clark St & Montrose Ave" "Broadway & Montrose Ave" "Clark St & Leland Ave" ...
##  $ start_station_id  : num [1:426887] 239 234 296 51 66 212 96 96 212 38 ...
##  $ end_station_name  : chr [1:426887] "Clark St & Leland Ave" "Southport Ave & Irving Park Rd" "Wilton Ave & Irving Park Rd" "Clark St & Leland Ave" ...
```

```
## $ end_station_id      : num [1:426887] 326 318 117 24 212 96 212 212 96 100 ...
## $ start_lat           : num [1:426887] 42 42 41.9 41.9 41.9 ...
## $ start_lng           : num [1:426887] -87.7 -87.7 -87.6 -87.6 -87.6 ...
## $ end_lat             : num [1:426887] 42 42 41.9 41.9 41.9 ...
## $ end_lng             : num [1:426887] -87.7 -87.7 -87.7 -87.6 -87.6 ...
## $ member_casual       : chr [1:426887] "member" "member" "member" "member" ...
## - attr(*, "spec")=
## .. cols(
## ..   ride_id = col_character(),
## ..   rideable_type = col_character(),
## ..   started_at = col_character(),
## ..   ended_at = col_character(),
## ..   start_station_name = col_character(),
## ..   start_station_id = col_double(),
## ..   end_station_name = col_character(),
## ..   end_station_id = col_double(),
## ..   start_lat = col_double(),
## ..   start_lng = col_double(),
## ..   end_lat = col_double(),
## ..   end_lng = col_double(),
## ..   member_casual = col_character()
## .. )
## - attr(*, "problems")=<externalptr>
```

Convert `ride_id` and `rideable_type` to character so that they can stack correctly

```
q1_2019 <- mutate(q1_2019, ride_id = as.character(ride_id),
,rideable_type = as.character(rideable_type))
```

Stack individual quarter's data frames into one big data frame

```
all_trips <- bind_rows(q1_2019, q1_2020)#, q3_2019)#, q4_2019, q1_2020)
```

Remove lat, long, birthyear, and gender fields as this data was dropped beginning in 2020

```
all_trips <- all_trips %>%
select(-c(start_lat, start_lng, end_lat, end_lng, birthyear, gender, "tripduration"))
```

```
#=====
```

**STEP 3: CLEAN UP AND ADD DATA TO PREPARE FOR ANALYSIS**

```
#=====
```

## Inspect the new table that has been created

```
colnames(all_trips) #List of column names
```

```
## [1] "ride_id"          "started_at"        "ended_at"
## [4] "rideable_type"     "start_station_id"  "start_station_name"
## [7] "end_station_id"    "end_station_name"  "member_casual"
```

```
nrow(all_trips) #How many rows are in data frame?
```

```
## [1] 791956
```

```
dim(all_trips) #Dimensions of the data frame?
```

```
## [1] 791956      9
```

```
head(all_trips) #See the first 6 rows of data frame. Also tail(all_trips)
```

```
## # A tibble: 6 x 9
##   ride_id started_at ended_at rideable_type start_station_id start_station_name
##   <chr>    <chr>      <chr>    <chr>                <dbl> <chr>
## 1 21742443 01-01-201~ 01-01-2~ 2167                199 Wabash Ave & Gran~
## 2 21742444 01-01-201~ 01-01-2~ 4386                44 State St & Randol~
## 3 21742445 01-01-201~ 01-01-2~ 1524                15 Racine Ave & 18th~
## 4 21742446 01-01-201~ 01-01-2~ 252                 123 California Ave & ~
## 5 21742447 01-01-201~ 01-01-2~ 1170                173 Mies van der Rohe~
## 6 21742448 01-01-201~ 01-01-2~ 2437                98 LaSalle St & Wash~
## # i 3 more variables: end_station_id <dbl>, end_station_name <chr>,
## #   member_casual <chr>
```

```
str(all_trips) #See list of columns and data types (numeric, character, etc)
```

```
## tibble [791,956 x 9] (S3: tbl_df/tbl/data.frame)
## $ ride_id      : chr [1:791956] "21742443" "21742444" "21742445" "21742446" ...
## $ started_at   : chr [1:791956] "01-01-2019 00:04" "01-01-2019 00:08" "01-01-2019 00:13" "01-01-2019 00:15" ...
## $ ended_at     : chr [1:791956] "01-01-2019 00:11" "01-01-2019 00:15" "01-01-2019 00:27" "01-01-2019 00:29" ...
## $ rideable_type : chr [1:791956] "2167" "4386" "1524" "252" ...
## $ start_station_id : num [1:791956] 199 44 15 123 173 98 98 211 150 268 ...
## $ start_station_name: chr [1:791956] "Wabash Ave & Grand Ave" "State St & Randolph St" "Racine Ave & 18th St" "California Ave & Washington St" ...
## $ end_station_id   : num [1:791956] 84 624 644 176 35 49 49 142 148 141 ...
## $ end_station_name : chr [1:791956] "Milwaukee Ave & Grand Ave" "Dearborn St & Van Buren St (*)" "Washington St & Racine Ave" "Washington St & Racine Ave" ...
## $ member_casual    : chr [1:791956] "Subscriber" "Subscriber" "Subscriber" "Subscriber" ...
```

```
summary(all_trips) #Statistical summary of data. Mainly for numerics
```

```
##   ride_id      started_at      ended_at      rideable_type
## Length:791956 Length:791956 Length:791956 Length:791956
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
##
##
##
## start_station_id start_station_name end_station_id end_station_name
## Min. : 2.0 Length:791956 Min. : 2.0 Length:791956
## 1st Qu.: 77.0 Class :character 1st Qu.: 77.0 Class :character
## Median :174.0 Mode :character Median :174.0 Mode :character
```

```
## Mean      :204.4      Mean      :204.4
## 3rd Qu.   :291.0      3rd Qu.   :291.0
## Max.      :675.0      Max.      :675.0
##          NA's       :1
## member_casual
## Length:791956
## Class :character
## Mode  :character
##
##
##
##
```

There are a few problems we will need to fix:

(1) In the “member\_casual” column, there are two names for members (“member” and

“Subscriber”) and two names for casual riders (“Customer” and “casual”). We will need to consolidate that from four to two labels. # (2) The data can only be aggregated at the ride-level, which is too granular. We will want to add some additional columns of data – such as day, month, year – that provide additional opportunities to aggregate the data. # (3) We will want to add a calculated field for length of ride since the 2020Q1 data did not have the “tripduration” column. We will add “ride\_length” to the entire dataframe for consistency. # (4) There are some rides where tripduration shows up as negative, including several hundred rides where Divvy took bikes out of circulation for Quality Control reasons. We will want to delete these rides. # In the “member\_casual” column, replace “Subscriber” with “member” and “Customer” with “casual” # Before 2020, Divvy used different labels for these two types of riders ... we will want to make our dataframe consistent with their current nomenclature

**N.B.: “Level” is a special property of a column that is retained even if a subset does not**

contain any values from a specific level # Begin by seeing how many observations fall under each usertype

```
table(all_trips$member_casual)
```

```
##
##      casual      Customer      member Subscriber
##      48480        23163      378407      341906
```

Reassign to the desired values (we will go with the current 2020 labels)

```
all_trips <- all_trips %>%
mutate(member_casual = recode(member_casual
, "Subscriber" = "member"
, "Customer" = "casual"))
```

Check to make sure the proper number of observations were reassigned

```
table(all_trips$member_casual)
```

```
##
## casual member
## 71643 720313
```

Add columns that list the date, month, day, and year of each ride

This will allow us to aggregate ride data for each month, day, or year ... before completing

these operations we could only aggregate at the ride level # <https://www.statmethods.net/input/dates.html>  
more on date formats in R found at that link

```
all_trips$date <- as.Date(all_trips$started_at) #The default format is yyyy-mm-dd
all_trips$month <- format(as.Date(all_trips$date), "%m")
all_trips$day <- format(as.Date(all_trips$date), "%d")
all_trips$year <- format(as.Date(all_trips$date), "%Y")
all_trips$day_of_week <- format(as.Date(all_trips$date), "%A")
```

Add a “ride\_length” calculation to all\_trips (in seconds)

<https://stat.ethz.ch/R-manual/R-devel/library/base/html/difftime.html>

```
all_trips$ride_length <- difftime(all_trips$ended_at, all_trips$started_at)
```

Inspect the structure of the columns

```
str(all_trips)
```

```
## tibble [791,956 x 15] (S3: tbl_df/tbl/data.frame)
## $ ride_id      : chr [1:791956] "21742443" "21742444" "21742445" "21742446" ...
## $ started_at   : chr [1:791956] "01-01-2019 00:04" "01-01-2019 00:08" "01-01-2019 00:13" "01-01-2019 00:15" ...
## $ ended_at     : chr [1:791956] "01-01-2019 00:11" "01-01-2019 00:15" "01-01-2019 00:27" "01-01-2019 00:29" ...
## $ rideable_type : chr [1:791956] "2167" "4386" "1524" "252" ...
## $ start_station_id : num [1:791956] 199 44 15 123 173 98 98 211 150 268 ...
## $ start_station_name: chr [1:791956] "Wabash Ave & Grand Ave" "State St & Randolph St" "Racine Ave & Grand Ave" ...
## $ end_station_id   : num [1:791956] 84 624 644 176 35 49 49 142 148 141 ...
## $ end_station_name : chr [1:791956] "Milwaukee Ave & Grand Ave" "Dearborn St & Van Buren St (*)" "Racine Ave & Grand Ave" ...
## $ member_casual    : chr [1:791956] "member" "member" "member" "member" ...
## $ date             : Date [1:791956], format: "1-01-20" "1-01-20" ...
## $ month            : chr [1:791956] "01" "01" "01" "01" ...
## $ day              : chr [1:791956] "20" "20" "20" "20" ...
## $ year             : chr [1:791956] "1" "1" "1" "1" ...
## $ day_of_week       : chr [1:791956] "Saturday" "Saturday" "Saturday" "Saturday" ...
```

```
## $ ride_length      : 'difftime' num [1:791956] 0 0 0 0 ...
##   ..- attr(*, "units")= chr "secs"
```

Convert “ride\_length” from Factor to numeric so we can run calculations on the data

```
is.factor(all_trips$ride_length)

## [1] FALSE

all_trips$ride_length <- as.numeric(as.character(all_trips$ride_length))
is.numeric(all_trips$ride_length)

## [1] TRUE
```

Remove “bad” data

The dataframe includes a few hundred entries when bikes were taken out of docks and

checked for quality by Divvy or ride\_length was negative # We will create a new version of the dataframe (v2) since data is being removed # <https://www.datasciencemadesimple.com/delete-or-drop-rows-in-r-with-conditions-2/>

```
all_trips_v2 <- all_trips[!(all_trips$start_station_name == "HQ QR" | all_trips$ride_length<0),]

#=====
```

## STEP 4: CONDUCT DESCRIPTIVE ANALYSIS

```
#=====
```

Descriptive analysis on ride\_length (all figures in seconds)

```
mean(all_trips_v2$ride_length) #straight average (total ride length / rides)

## [1] 114170.8

median(all_trips_v2$ride_length) #midpoint number in the ascending array of ride lengths

## [1] 0

max(all_trips_v2$ride_length) #longest ride

## [1] 946684800

min(all_trips_v2$ride_length) #shortest ride

## [1] 0
```



You can condense the four lines above to one line using `summary()` on the specific attribute

```
summary(all_trips_v2$ride_length)

##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
##         0         0         0    114171         0 946684800
```

Compare members and casual users

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual, FUN = mean)
```

```
## all_trips_v2$member_casual all_trips_v2$ride_length
## 1          casual          696019.11
## 2          member          59395.69
```

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual, FUN = median)
```

```
## all_trips_v2$member_casual all_trips_v2$ride_length
## 1          casual          0
## 2          member          0
```

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual, FUN = max)
```

```
## all_trips_v2$member_casual all_trips_v2$ride_length
## 1          casual    946684800
## 2          member    820454400
```

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual, FUN = min)
```

```
## all_trips_v2$member_casual all_trips_v2$ride_length
## 1          casual          0
## 2          member          0
```

See the average ride time by each day for members vs casual users

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual + all_trips_v2$day_of_week,
FUN = mean)
```

```
## all_trips_v2$member_casual all_trips_v2$day_of_week all_trips_v2$ride_length
## 1          casual          Friday          742739.29
## 2          member          Friday          68795.18
## 3          casual          Monday          485311.28
## 4          member          Monday          53084.27
## 5          casual          Saturday          849717.58
## 6          member          Saturday          55277.17
## 7          casual          Sunday          983708.54
## 8          member          Sunday          45735.78
## 9          casual          Thursday          554798.24
## 10         member          Thursday          56110.19
## 11         casual          Tuesday          758923.40
## 12         member          Tuesday          56182.94
## 13         casual          Wednesday          680982.83
## 14         member          Wednesday          83657.82
```

Notice that the days of the week are out of order. Let's fix that.

```
all_trips_v2$day_of_week <- ordered(all_trips_v2$day_of_week, levels=c("Sunday", "Monday",  
"Tuesday", "Wednesday", "Thursday", "Friday", "Saturday"))
```

Now, let's run the average ride time by each day for members vs casual users

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual + all_trips_v2$day_of_week,  
FUN = mean)
```

```
##      all_trips_v2$member_casual all_trips_v2$day_of_week all_trips_v2$ride_length  
## 1                casual      Sunday      983708.54  
## 2                member      Sunday      45735.78  
## 3                casual      Monday      485311.28  
## 4                member      Monday      53084.27  
## 5                casual      Tuesday      758923.40  
## 6                member      Tuesday      56182.94  
## 7                casual      Wednesday      680982.83  
## 8                member      Wednesday      83657.82  
## 9                casual      Thursday      554798.24  
## 10               member      Thursday      56110.19  
## 11               casual      Friday      742739.29  
## 12               member      Friday      68795.18  
## 13               casual      Saturday      849717.58  
## 14               member      Saturday      55277.17
```

```
install.packages("lubridate") library(lubridate) install.packages("dplyr") # For only dplyr library(dplyr)
```

OR (for tidyverse, which includes dplyr and other useful packages)

```
install.packages("tidyverse") library(tidyverse) all_trips_v2 <- all_trips_v2 %>% mutate(started_at =  
ymd_hms(started_at)) # Convert started_at to POSIXct format
```

analyze ridership data by type and weekday

```
library(lubridate)  
library(dplyr)  
  
all_trips_v2 %>%  
  mutate(  
    started_at = ymd_hms(started_at), # Ensure 'started_at' is in proper date-time format  
    weekday = wday(started_at, label = TRUE) # Create weekday field  
  ) %>%  
  group_by(member_casual, weekday) %>% # Group by user type and weekday  
  summarise(  
    number_of_rides = n(), # Calculate number of rides  
    average_duration = mean(ride_length, na.rm = TRUE) # Calculate average ride length  
  ) %>%  
  arrange(member_casual, weekday) # Sort by user type and weekday
```

```
## `summarise()` has grouped output by 'member_casual'. You can override using the
## `.groups` argument.
```


```
## # A tibble: 14 x 4
## # Groups:   member_casual [2]
##   member_casual weekday number_of_rides average_duration
##   <chr>          <ord>          <int>          <dbl>
## 1 casual        Sun              8119          983709.
## 2 casual        Mon             11510          485311.
## 3 casual        Tue             13041          758923.
## 4 casual        Wed              8913          680983.
## 5 casual        Thu             13668          554798.
## 6 casual        Fri              5574          742739.
## 7 casual        Sat              6981          849718.
## 8 member        Sun            113887           45736.
## 9 member        Mon            93325           53084.
## 10 member       Tue            116923           56183.
## 11 member       Wed            102747           83658.
## 12 member       Thu            105723           56110.
## 13 member       Fri             84904           68795.
## 14 member       Sat            102760           55277.
```

Let's visualize the number of rides by rider type

```
all_trips_v2 %>%
mutate(weekday = wday(started_at, label = TRUE)) %>%
group_by(member_casual, weekday) %>%

summarise(number_of_rides = n()
,average_duration = mean(ride_length)) %>%
arrange(member_casual, weekday) %>%
ggplot(aes(x = weekday, y = number_of_rides, fill = member_casual)) +
geom_col(position = "dodge")
```

```
## `summarise()` has grouped output by 'member_casual'. You can override using the
## `.groups` argument.
```



Let's create a visualization for average duration

```
all_trips_v2 %>%
  mutate(weekday = wday(started_at, label = TRUE)) %>%
  group_by(member_casual, weekday) %>%
  summarise(number_of_rides = n()
    ,average_duration = mean(ride_length)) %>%
  arrange(member_casual, weekday) %>%
  ggplot(aes(x = weekday, y = average_duration, fill = member_casual)) +
  geom_col(position = "dodge")
```

## `summarise()` has grouped output by 'member\_casual'. You can override using the  
## `.groups` argument.

